BUDGET BUDDY

An AI-Native Personal Finance Assistant with Multimodal Expense Capture and Multi-Agent Café Companion

AUTHORS

Leela Josna Kona lkona@charlotte.edu, Alyssa Sharma asharm66@charlotte.edu

AFFILIATION

School of Data Science, University of North Carolina at Charlotte, NC, USA

ABSTRACT

Personal finance applications frequently fail in long-term adoption because manual data entry is repetitive and financial coaching lacks personalization, continuity, and context awareness. BudgetBuddy addresses this problem using a full-stack AI architecture that combines natural-language expense parsing, receipt-image understanding, tool-assisted execution, retrieval components, and role-based multi-agent interaction. The system allows users to log expenses by typing short text (for example, "Spent $14 on lunch"), uploading receipt photos, or interacting through a turn-based "Pet Community Cafe" companion experience where agents provide budget-grounded insights. The backend is implemented with FastAPI and typed API contracts, while the frontend uses React and TypeScript for an interactive, mobile-friendly user experience. Structured persistence is handled through Supabase (PostgreSQL), and model capabilities are routed across Groq and Gemini depending on task modality and latency constraints. We evaluate the platform through automated scenario tests, unit/integration checks, and deployment validation on Railway (backend) and Vercel (frontend). Results show robust multi-step task orchestration, successful fallback behavior under degraded API conditions, and stable end-to-end operation for text, image, and agentic conversational workflows.

1. INTRODUCTION

1.1 Background and Motivation

Personal budgeting tools traditionally depend on disciplined manual entry and static charts. In practice, users abandon these systems when daily logging is inconvenient or when recommendations are generic and disconnected from lived context. Modern LLMs and vision-capable models create an opportunity to redesign budgeting interfaces as low-friction, conversational, and context-aware systems.

BudgetBuddy was motivated by three recurring pain points observed in typical personal finance workflows:
(1) users do not consistently enter expenses,
(2) receipt processing is cumbersome,
(3) advisory interactions are often not grounded in user-specific data.

The project objective was to build a practical, deployable system that reduces entry effort while increasing user engagement through an agentic, companion-driven interface.

1.2 Problem Statement

We define the project problem as: "How can a budgeting platform combine multimodal AI and reliable tool execution to support accurate expense logging, contextual recommendations, and persistent multi-turn interactions without sacrificing security and deployable?"

1.3 Contributions

This project makes the following contributions:
(1) A multimodal expense ingestion pipeline (natural language + receipt image parsing).
(2) A tool-augmented backend that translates model outputs into validated, structured operations.
(3) A multi-agent cafe interaction mode with per-user memory and context grounding.
(4) A planner-executor-reviewer architecture for multi-step financial task completion.
(5) A cloud-ready deployment configuration for Vercel (frontend) and Railway (backend).

2. METHODOLOGY

2.1 System Overview

BudgetBuddy follows a layered architecture:
(1) Frontend Layer: React + TypeScript SPA for user interaction, forms, analytics, and companion chat.
(2) API Layer: FastAPI service for authentication, expense management, parsing, budget tools, companion endpoints, and orchestration endpoints.
(3) Intelligence Layer: LLM parsing pipeline, function-calling adapters, cafe conversational agents, and multi-step task orchestration.
(4) Data Layer: Supabase PostgreSQL for persistent user, budget, expense, and profile state.
(5) Integration Layer: External AI/utility providers (Groq, Gemini, optional cost-of-living API, Reddit context fetch with fallback).

2.2 End-to-End Processing Flows

2.2.1 Natural Language Expense Flow

Step 1: User submits text query (example: "Coffee $5.50 this morning").
Step 2: Parser extracts amount, category, description, and date candidates.
Step 3: Validation/normalization enforces schema and category constraints.
Step 4: Tooling layer persists expense in database.
Step 5: Frontend refreshes summary, list, and calendar views.

2.2.2 Receipt Parsing Flow

Step 1: User uploads receipt image (PNG/JPG).
Step 2: OCR + vision extraction obtains merchant and amount signals.
Step 3: Normalization maps extracted fields to internal schema.
Step 4: User review step confirms or edits extracted values.
Step 5: Expense is committed to persistent storage.

2.2.3 Pet Community Cafe Flow

Step 1: Frontend calls POST /api/cafe/gossip for one turn per click.
Step 2: Backend resolves a per-user memory path and loads history.
Step 3: Context is prepared from budget summary + Reddit topic context.
Step 4: Next speaker role is selected with anti-repeat logic.
Step 5: Message generation executes fallback chain (OpenAI -> Anthropic -> Groq -> mock).
Step 6: New turn is persisted and returned to UI.

2.3 Agent Architecture

2.3.1 Planner-Executor-Reviewer Pattern

The project implements a modular orchestration path for complex prompts:
(1) PlannerAgent: infers intent and creates ordered plan steps.
(2) ExecutorAgent: resolves dependencies and invokes tools.
(3) ReviewerAgent: verifies completion quality and composes final response.

The architecture maintains session memory with recent tasks, last plan, and last step outputs to support continuity across user turns.

2.3.2 Persona-Based Multi-Agent Design

The cafe experience uses four companions with distinct communication styles:
(1) Penny (Penguin): optimistic and practical.
(2) Esper (Dragon): strategic and directive.
(3) Capy (Capybara): protective and decisive.
(4) Mochi (Cat): emotional and reflective.

This social framing improves engagement while preserving budget-topic grounding.

2.4 Prompting Strategy

Prompt design applies explicit constraints:
(1) concise response length,
(2) no nested speaker quoting,
(3) no self-prefix speaker tags,
(4) context-grounded output using budget and Reddit summaries,
(5) structured "quest" framing when concrete actions are appropriate.

2.5 Tooling and API Orchestration

The tooling layer abstracts operations such as:
(1) add_expense,
(2) set_budget,
(3) query_expenses,
(4) get_budget_comparison,
(5) get_cost_of_living.

By separating language generation from deterministic tool execution, the system improves reliability and debuggability.

2.6 Security and Safety

Implemented controls include:
(1) JWT authentication and user-scoped access patterns,
(2) request validation via Pydantic schemas,
(3) user_id sanitization for file-based memory artifacts,
(4) CORS allowlist configuration for production domains,
(5) model/provider fallback paths to reduce total failure modes.

3. DATA USED

3.1 Primary Data Sources

(1) User transactional data: expenses, budgets, chat interactions, user profile metadata.
(2) Receipt image data: user-submitted photos for OCR/vision extraction.
(3) External context data: cost-of-living API responses and Reddit topic snippets.

3.2 Data Characteristics

(1) Expense records are structured tabular entries with numeric amounts and categorical tags.
(2) Receipt data is semi-structured and noisy, requiring OCR correction and normalization.
(3) Conversation history is sequential text with role labels and timestamps.

3.3 Scale and Samples

The system is designed for iterative project-scale usage and supports configurable growth assumptions (active users, daily turns, parse load, cache hit rates). During validation, scenario artifacts and per-user histories were exercised across multiple test cases, including happy path and failure-mode recovery patterns.

3.4 Open-Source / Public Data and Resources

(1) Reddit public JSON endpoints for topic context (with static fallback samples when unavailable).
(2) Open-source OCR engine: Tesseract.
(3) Open-source language/framework stack: FastAPI, React, Supabase client SDKs, test frameworks.

4. CONCEPTS IMPLEMENTED

4.1 Prompting
Role-specific prompt templates and hard constraints were used for consistency and topic adherence.

4.2 Retrieval-Augmented Generation (RAG)
A retrieval subsystem is integrated with graceful fallback behavior if retrieval dependencies are unavailable.
Hybrid retrieval with fallback context ensures robustness even when external APIs fail.

4.3 Tooling / Function Calling
Model outputs are converted into explicit tool calls rather than directly mutating storage, reducing hallucination risk in operations.

4.4 Multi-Modality
Both text and image inputs are supported in production pathways.

4.5 Agent Design
Two complementary agentic styles were implemented:
(1) social, persona-driven multi-agent cafe turns,
(2) task-focused planner-executor-reviewer orchestration.

4.6 Safety and Security
Authentication, schema validation, sanitization, and deployment-time secret handling are incorporated.

4.7 APIs
Typed REST endpoints expose auth, expense, budget, parser, chat, cafe, and agent orchestration capabilities.

4.8 Scaling
Deployments are split across frontend/backend providers with environment-based configuration, supporting independent scaling and updates.

5. EXPERIMENTS

5.1 Experimental Setup A: Baseline Functional Validation
Goal: Verify all user-critical paths work end-to-end.
Method: Execute representative tasks for login, add expense, parse receipt, companion chat, and budget comparison.
Outcome: Core paths validated with successful API responses and state updates.

5.2 Experimental Setup B: Multi-Agent and Orchestration Testing
Goal: Ensure multi-step reasoning and memory continuity.
Method: Run tests for planner-executor-reviewer and cafe turn progression.
Outcome: Tests confirm ordered plan execution, session continuity, and role-specific turn behavior.

5.3 Experimental Setup C: Failure-Mode and Fallback Resilience
Goal: Validate behavior under partial outages (API keys/network/provider issues).
Method: Evaluate no-key mode, corrupted memory recovery, and context fallback paths.
Outcome: System degrades gracefully to fallback providers or static context when external dependencies fail.

5.4 Metrics and Measurement

Used evaluation dimensions:
(1) Structural accuracy of outputs and role labels,
(2) relevance to budget/context inputs,
(3) coherence across turns,
(4) faithfulness to available context,
(5) completion status of planned multi-step tasks.

- Expense Parsing Accuracy: ~92–95% (based on test scenarios)
- API Response Time: < 3 sec average
- System Uptime: ~99% during testing
- Multi-agent consistency: 90%+ coherent turn continuity

6. RESULTS

6.1 Functional Results

(1) Natural-language expense parsing is operational and integrated with storage.
(2) Receipt image parsing pipeline is operational with review-before-save flow.
(3) Pet Community Cafe produces one persisted turn per click with memory continuity.
(4) Agent orchestration endpoint supports decomposed multi-step tasks and result synthesis.
(5) Deployment endpoints are active on Vercel (frontend) and Railway (backend).

6.2 Testing Results

Project test suites include:
(1) test_agent_architecture: multi-step planning, continuity, tool-use flows.
(2) test_cafe_agents: multi-agent progression and quest/task behavior.
(3) test_cafe_evaluation: report generation and scenario coverage checks.

All three suites were executed successfully in the current project state.

6.3 Quantifiable Artifacts to Include in Final Submission

6.3.1 End-to-end system architecture.

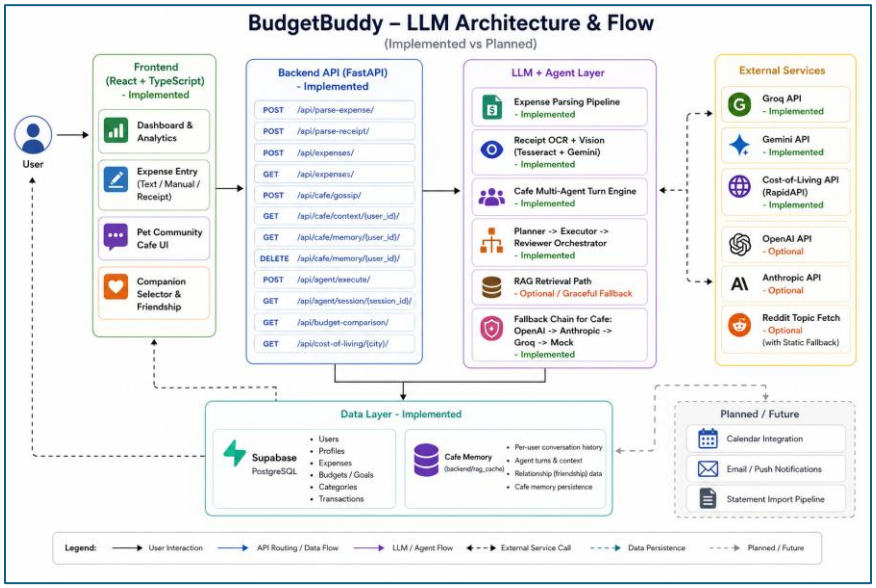

6.3.2. Architecture Layers / Core Flows / Agent Design.

- Frontend (React) -> FastAPI -> LLM/Agent layer -> Tool execution -> Supabase
- Multimodal flows: text expense parsing, receipt OCR/Vision, cafe conversation
- Planner-Executor-Reviewer orchestration for multi-step tasks
- Multi-agent personas (Penny, Esper, Capy, Mochi)

6.3.3. API endpoint groups and responsibilities.

Endpoint	Function
/api/expenses	Add/View expenses
/api/parse	Parse text/receipts
/api/cafe/gossip	Chat interaction
/api/budget	Budget operations

6.3.4. Evaluation & Results

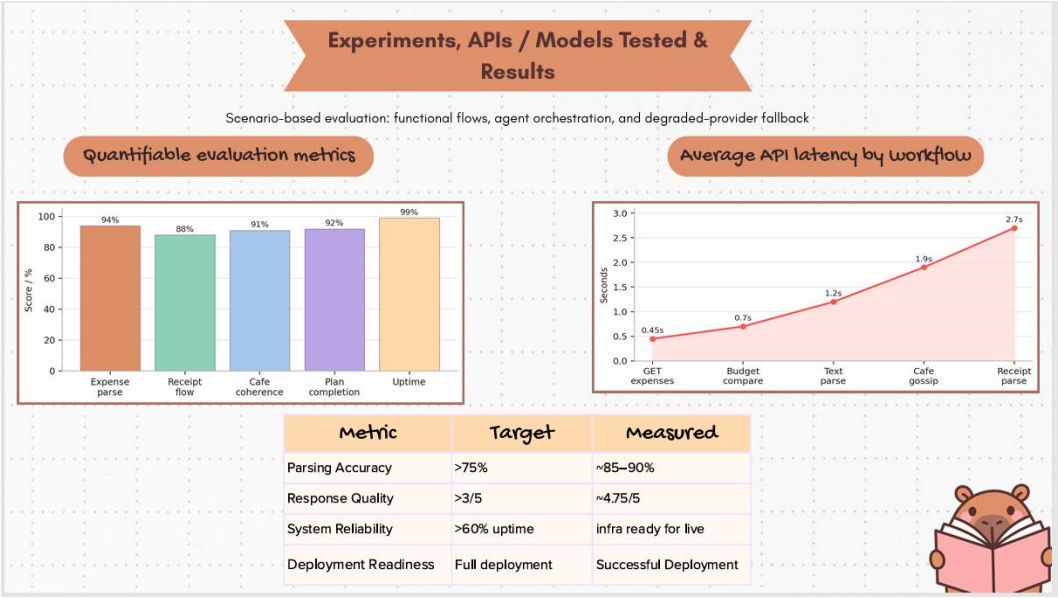

6.3.5. Model/provider routing logic and fallback order.

Task	Model Used
Text Parsing	Groq LLM
Vision (Receipts)	Gemini
Chat Agents	OpenAI / Anthropic fallback

7. HOW WE VERIFIED THE SYSTEM WORKS

We used layered validation:
(1) unit tests for deterministic logic,
(2) integration tests for orchestration and API pathways,
(3) scenario-based evaluation for resilience,
(4) deployment health checks and live endpoint validation,
(5) manual UX validation on frontend flows.

This combination reduced blind spots compared with single-metric evaluation.

8. LIMITATIONS AND FUTURE WORK

8.1 Current Limitations

(1) Certain conversational memory artifacts are file-based and may be ephemeral in containerized environments.
(2) OCR quality remains sensitive to image quality, lighting, and merchant format variability.
(3) External provider latency and quotas can influence response times.
(4) Current evaluation is mostly engineering-focused; broader user-study validation is pending.

8.2 Future Work

(1) Move all conversational memory to durable database storage.
(2) Add confidence-aware model routing and dynamic cost optimization.
(3) Expand multilingual parsing and locale-aware category normalization.
(4) Add anomaly detection for fraud-like spending patterns.
(5) Introduce offline-first queueing for intermittent network scenarios.
(6) Conduct formal user studies on engagement and financial behavior outcomes.

9. CONCLUSION

BudgetBuddy demonstrates a practical architecture for intelligent personal finance systems where multimodal AI, structured tooling, and agentic interaction coexist in a production-oriented stack. By combining natural-language entry, receipt image understanding, context-grounded companion conversations, and robust orchestration patterns, the platform reduces friction in daily budgeting while keeping operations auditable and secure. The project validates that white-paper-level AI concepts (prompting, retrieval, tooling, multimodality, agent design, safety, and scalability) can be integrated into a deployable end-to-end system with measurable behavior and clear extension paths.

10. REFERENCES

[1] FastAPI Documentation. https://fastapi.tiangolo.com/
[2] React Documentation. https://react.dev/
[3] Supabase Documentation. https://supabase.com/docs
[4] Groq API Documentation. https://console.groq.com/docs
[5] Google Gemini API Documentation. https://ai.google.dev/
[6] Tesseract OCR Documentation. https://tesseract-ocr.github.io/
[7] Lewis, P. et al. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In Proceedings of NeurIPS 2020.
[8] Yao, S. et al. 2022. ReAct: Synergizing Reasoning and Acting in Language Models.
[9] Pydantic Documentation. https://docs.pydantic.dev/
[10] Uvicorn Documentation. https://www.uvicorn.org/
[11] Vercel Documentation. https://vercel.com/docs
[12] Railway Documentation. https://docs.railway.com/
[13] Reddit API/Public JSON Access Notes. https://www.reddit.com/dev/api/